

Speaker scripts

Spoken lines only, in running order — the rehearsal copy. Everything else (headlines, exhibits, backup numbers) lives in the deck brief.

1137 words across 10 slides · **8:45** at a measured 130 wpm · **8:07** at a normal 140 wpm · the per-slide budgets sum to 10:10. Hard cut-off is 10:00; target finishing by 9:00.

Generated by `scripts/export_scripts_and_qa.py` from
`docs/presentation/deck_content.md`. Edit the deck brief, not this file.

Slide 1 · Title

Tuan · budget ~30s · 58 words (~0:27 at 130 wpm)

Good morning — we're Cultural Blend: I'm Tuan, and my teammate Sophia will take you through what this means for the business. Kiva is built on stories — every loan page leads with one, the way a landing page leads with copy. We asked 1.45 million real loans a single question: does the story actually move the money?

Slide 2 · The question

Tuan · budget ~50s · 81 words (~0:37 at 130 wpm)

Why should a marketing audience care? Because on Kiva, the lender is the customer and the loan page is the product page. Loan size, sector, geography — fixed at listing. The story is the one element a platform can coach, test, and optimise — classic conversion territory. So: when a story leans on family, competence, or urgency, does the loan fund faster? And does that depend on who's asking, and when? That's a testable claim — so we tested it, hard.

Slide 3 · How we stress-tested our findings

Tuan · budget ~65s · 110 words (~0:51 at 130 wpm)

Two disciplines before any findings — because in a dataset this size it is dangerously easy to find things that aren't there. For prediction: train only on the past, score only on loans posted in 2024–25 — the models never see the future they're graded on. For the framing claims: every 'significant' result had to survive country-clustered standard errors — ten thousand loans from one country are not ten thousand independent customers — and where a result rested on just a couple of countries, an even harsher few-cluster screen on top. Most headline-looking results did not survive. Hold that thought while Sophia shows you what the market itself looks like.

Slide 4 · A marketplace that hasn't recovered

Sophia · budget ~65s · 111 words (~0:51 at 130 wpm)

Thanks, Tuan. Start with the market, because every other number lives inside it. Before the pandemic, 46% of loans funded within 24 hours — nearly half converting same-day. Since 2020 it's been under a third — and through 2025, the 'recovery' is minus 0.3 points. With over half a million loans on each side of that divide, this isn't noise: it's a structurally slower, more selective marketplace. For a marketer that reframes the whole job — when customers get pickier, knowing what actually drives conversion matters more, not less. Four years without a recovery is also why our recommendations start with acting rather than waiting for the market to fix itself.

Slide 5 • Structure beats story

Sophia • budget ~65s • 98 words (~0:45 at 130 wpm)

And what drives it is structure. Loan amount and repayment terms correlate with funding speed roughly ten times more strongly than any narrative signal. Sector alone spans an order of magnitude — sanitation stories fund in under a day, clothing takes twelve. And the starkest gap in the data: loans posted by women fund in a median of 2.3 days, by men 7.7 — three times longer. None of this is causal, but it's enormous, it's stable, and it dwarfs anything the words do. Which brings us to the question we actually came here to answer. Tuan —

Slide 6 • What survives scrutiny

Tuan • budget ~1m50s • 200 words (~1:32 at 130 wpm)

So: does the story matter? Our honest answer: no narrative result is robust enough across specifications to support a recommendation. Urgency language looked like a universal win — significant at p below 0.001. Cluster by country, and it collapses to 0.44. Gone. Family framing — and here our own first version got it wrong: we tested whether regions differ from Africa, which is not the same as whether family framing helps within a region. Corrected, two pooled categories — Palestine plus Yemen, and Honduras plus Nicaragua — do show faster funding in every fit we ran. But each rests on exactly two countries, and against a deliberately harsh few-cluster screen — a **conservative heuristic**, not calibrated inference: a t distribution with one degree of freedom, critical value **12.7, not 1.96** — neither is significant: p between 0.06 and 0.21. So we report a hypothesis worth testing, not a finding. Sentiment is the honest illustration: it survives country clustering in one of our two specifications and not the other — genuinely open, robust in neither direction. We'd rather report **no robust evidence** than one exciting result we can't defend — because a recommendation you'd ship to real borrowers deserves that bar.

Slide 7 • Beyond keywords

Tuan • budget ~50s • 101 words (~0:47 at 130 wpm)

One more layer before the recommendations — we went past keyword counting. Topic modelling finds eight coherent story themes, and mean funding speed swings ninefold across them — descriptively, not as a tested effect: sanitation and clean-water stories in under two days, group farming closer to two weeks. But notice what a theme mostly encodes: what the loan is FOR. Structure again — not persuasion — and it's why the structural review we're about to recommend should cover what loans are FOR, not just which sector they sit in. So what should Kiva actually do with all of this? Sophia —

Slide 8 · What this means in practice

Sophia · budget ~2m · 264 words (~2:02 at 130 wpm)

Three moves for the platform, one for borrowers — and notice the verbs, because the evidence sets them. One: don't ship writing tips — a firm don't, because a robust null is the one thing we did find. For the content team that's a build spared; for borrowers, it's not being coached into copy we found **no robust evidence** for — and the timing advice belongs to field partners, not borrowers: a bigger ask sits on the platform longer — two days for the smallest loans against nineteen for the largest — and since **96% of loans** in this data are already disbursed before the page even goes live, that is a **partner capital-cycle** fact, not a borrower's waiting time. Which is also why we are not telling anyone to ask for less. Two — an experiment, not an action, because there we have a pattern we could not confirm: for the growth team and the field partners in exactly four markets — Palestine, Yemen, Honduras, Nicaragua — run the **country-stratified A/B test**: family-framing prompt versus standard at listing. **test before you ship**. Three — review, not change, because these are associations and not causes: for the product team, the structural gaps are the real levers — how the consistently slower sectors and regions get surfaced and supported — because those gaps are ten times any wording effect. And the classifier stays a prototype, not a deployment — a retrospective ranking prototype among funded loans — **AUC 0.90** on strictly future data — so data science retrains on all listings, including expired ones, first.

Slide 9 · What this can't tell us

Sophia · budget ~35s · 63 words (~0:29 at 130 wpm)

Three honest limits. Association, never causation — nobody randomly assigned writing styles. We measure how fast funded loans fund — not whether a loan funds at all. And our framing measures are transparent, simple rules — not every nuance of persuasion. We'd rather you know exactly what this can and cannot say — that's what makes the parts we do claim worth trusting.

Slide 10 · Closing

Both · budget ~20s · 51 words (~0:24 at 130 wpm)

(Sophia) In this data, the story barely registers — the structure carries the signal. (Tuan) And testing hard enough to say, honestly, that there is **no robust evidence** for the story — that's worth more to a platform than a good-sounding tip. Thank you — we're happy to take your questions.